

# Introduction to Parallel Computing

Jinzhen Wang  
jwang96@charlotte.edu

Department of Computer Science  
UNC Charlotte

# Learning Outcome

At this end of this lecture, you will be able to

- Measure the performance of parallelism.
- Articulate why not all problems will benefit from parallelism.
- Articulate the difference between weak and strong scaling.
- Articulate the difference between horizontal and vertical scaling.

# Why not just measure time?

The end goal of parallelism is to decrease the time of some calculation.  
One can measure time and report time improvements.

- Often raw time is hard to make sense of.
- Not all applications use time as the metrics.
  - Video games: FPS
- Makes it hard to get insight across different machines.

# Measure of Success: speedup

## Definition

$S(N) = \frac{T_1}{T_N}$  where  $T_1$  is sequential time and  $T_N$  is time on  $N$  processors.  
You should use as  $T_1$  the best sequential time you can get.

## Example

A sequential application taking 100 minutes, being executed in parallel on 4 processors to take 50 minutes has a speedup on 4 processors:

$$S(4) = \frac{100}{50} = 2$$

Speed is the key to measure how well your parallel computing works.

# What speedup is good?

## Linear speedup

Ideally, you would want to achieve linear speedup:  $S(N) = N$ .

If you achieve linear speedup (or close), the code is said to scale well.

It is ideal, you almost never get  $S(N) = N$

## Super linear speedup

You can sometimes get  $S(N) > N$ .

It is usually an anomaly that comes from the fact that the parallel algorithm does less computation than the sequential one in that particular case.

- if you double the population of a city, you increase the productivity of the city by more than double (roughly 115%).
- if a problem can be divided into pieces that fit entirely into a CPU's cache, a significant speedup can be observed.


# Why wouldn't all computation scale?

## Inherently sequential computation

- Two ovens won't cook a cake faster.
- Combining two half-cups of water at  $50^{\circ}\text{C}$  will not result in one cup of water at  $100^{\circ}\text{C}$ .

# Why wouldn't all computation scale?

## Communications/Synchronization



# How much can a computation scale ?

## Strong Scaling

In a strong scaling experiment, you pick one particular test case that you want to solve. And you see how fast that one can be solved.

- + The best kind of experiment
  - Scaling depends on problem size
  - How to pick a representative problem?



# How much can a computation scale ?

## Amdahl's Law


- A fraction  $x$  of the computation that is serial.
- The rest  $(1 - x)$  is completely parallel

Time on  $N$  computational units

$$T(N) \geq T(1)(x + \frac{1-x}{N})$$

Speedup on  $N$  computational units

$$S(N) \leq \frac{1}{x}$$



# How much can a computation scale?

## Weak Scaling


In a weak scaling experiment, you increase the volume of computation (in term of complexity) with the amount of resource. (Usually report time.)

- + Often easier to achieve scaling
- Not all problems have increasing size

# How much can a computation scale?

## Gustafson's Law

- The sequential part  $x$  is a one time fixed overhead **independent** of the amount of work
- The rest of the computation scales linearly
- Then you can process a dataset of size  $N$  with  $N$  processors in the same amount of time 1 processor would a dataset of size 1.
- $S(N) = x + N(1 - x)$ . with  $x$  the fraction the overhead represents on 1 computational unit.



# Examples of Strong/Weak Scaling?

- Strong Scaling
  - Decode a passcode
  - Image processing
- Weak Scaling
  - Weather forecast
  - Protein folding

# Scaling systems

## Metric: Latency

Time it takes to perform one computation.

## Metric: Bandwidth

Rate at which one can perform computation

# Scaling systems

## Metric: Latency

Time it takes to perform one computation.

## Vertical Scaling

Vertically scaling adds resources into a single node to increase its computational speed.

The goal is often to improve the latency of the system

## Metric: Bandwidth

Rate at which one can perform computation

## Horizontal Scaling

Horizontal scaling adds resources as additional nodes to increase the computational speed.

The goal is often to improve the bandwidth of the system

# Scaling systems

## Metric: Latency

Time it takes to perform one computation.

## Vertical Scaling

Vertically scaling adds resources into a single node to increase its computational speed. The goal is often to improve the latency of the system

- Increasing speed limits will decrease commute time.
- Using better servers to decrease query time.

## Metric: Bandwidth

Rate at which one can perform computation

## Horizontal Scaling

Horizontal scaling adds resources as additional nodes to increase the computational speed.

The goal is often to improve the bandwidth of the system

- Adding a lane to the highway will increase flow.
- Adding servers to process more clients.

# External

On parallel computing:

- Tim Mattson on why parallel computing: <https://www.youtube.com/watch?v=cMWGeJyrc9w>
- Tim Mattson on concurrency and parallelism: <https://www.youtube.com/watch?v=6jFkNjhJ-Z4>
- Wikipedia on parallel computing [https://en.wikipedia.org/wiki/Parallel\\_computing](https://en.wikipedia.org/wiki/Parallel_computing)
- Amdahl's Law <https://en.wikipedia.org/wiki/Amdahl's Law>
- "Parallel Programming in MPI and OpenMP" by Victor Eijkhout  
<https://bitbucket.org/VictorEijkhout/parallel-computing-book/raw/94518afb68da8f5af7104c0ce2343a8dc04f3a16/EijkhoutParComp.pdf>

Books that could be useful throughout the semester (most e-copy available through the library):

- Sushil K. Prasad, Anshul Gupta, Arnold Rosenberg, Alan Sussman, and Charles Weems. Topics in Parallel and Distributed Computing. Enhancing the Undergraduate Curriculum: Performance, Concurrency, and Programming on Modern Platforms. Springer 2018.  
<https://link.springer.com/book/10.1007/978-3-319-93109-8>
- Barbara Chapman, Gabriele Jost, and Ruud van der Pas. Using OpenMP. Portable Shared Memory Parallel Programming. MIT Press. 2007. Available through the library at <https://librarylink.uncc.edu/login?url=http://ieeexplore.ieee.org/servlet/opac?bknumber=6267237>
- Using MPI, 3rd edition. William Gropp, Ewing Lusk and Anthony Skjellum. MIT Press. Available through the library at <https://librarylink.uncc.edu/login?url=http://ieeexplore.ieee.org/xpl/bkabstractplus.jsp?bkn=6981847>