

Parallel Task Graphs

Jinzhen Wang
jwang96@charlotte.edu

Department of Computer Science
UNC Charlotte

Learning Outcomes

Lecture

At the end of this session you will know how to:


- Give two representations of parallel codes
- Compute metrics on parallel task graphs
- Interpret metrics of parallel task graphs in term of parallel execution

The Parallel Task Graph representation (PTG)

DAG representation

- Represents tasks as vertices.
- Represents x before y using a $x \rightarrow y$ directed edge.
- The graph is always without cycles.

Example




The Parallel Task Graph representation (PTG)

DAG representation


- Represents tasks as vertices.
- Represents x before y using a $x \rightarrow y$ directed edge.
- The graph is always without cycles.
- Processing time required associated with vertices, often denoted p_i

Example



A lemon pie recipe

- 1 break 2 eggs and split the white and yoke
- 2 cut 125g of butter in cubes
- 3 mix yoke and 70g of sugar+5cl of water
- 4 mix 250g of flour with butter
- 5 mix (3) and (4) and make a ball
- 6 spread (5)
- 7 heat oven to 180C
- 8 put crust (6) in pie pan
- 9 wash 4 lemons
- 10 peel two lemons from (9) and finely cut them
- 11 press two lemons from (9)
- 12 mix lemons(11), peel(10), 160g of sugar, 1 sp of flour
- 13 cook slowly (12)
- 14 whip 3 eggs
- 15 mix (14) and (13) and cook fast whipping
- 16 empty (15) in (17)
- 17 cook (8) for 20 minutes
- 18 wait until (16) cools



The conflict graph representation

Example

Conflict graph

- Used to represent a set of tasks that can be executed in any order but that use a common resource.
- Undirected graph with edges that connect tasks with a conflict.

Class	Instructor	Lab
I	A	1
II	B	1
III	A	2
IV	C	2
V	C	1


The conflict graph representation

Conflict graph

- Used to represent a set of tasks that can be executed in any order but that use a common resource.
- Undirected graph with edges that connect tasks with a conflict.

Example

Class	Instructor	Lab
I	A	1
II	B	1
III	A	2
IV	C	2
V	C	1



Conflict Graph vs. Dependency Graph

conflict graphs and **dependency graphs** are used in parallel computing to represent relationships between tasks or operations, but they serve different purposes and emphasize different aspects of these relationships.

Dependency Graph

- Purpose - shows the **order** where tasks or operations must be executed. It is used to visualize and analyze the dependencies between tasks, where one task must be completed before another can start.
- Nodes: Each node represents a task or operation.
- Edges: Directed edges ($A \rightarrow B$) represent dependencies, meaning task A must be completed before task B can start. The graph is usually DAG.
- Use Case: Crucial in **scheduling** tasks, optimizing execution order, and understanding parallelization limits. They help determine the **critical path**, which dictates the minimum time to complete all tasks.
- Example: In the robotcoin problem, a dependency graph would show how the computation of each cell in F depends on the cells to the left and above it.

Conflict Graph

- Purpose: represents **mutual exclusivity** in task execution, typically in the context of shared resources or parallel execution.
- Each node represents a task or operation.
- A (undirected) edge between two nodes (A - B) indicates that tasks A and B cannot be executed in parallel due to a conflict.
- Use case: Useful in scenarios like resource allocation, where tasks must be scheduled in such a way that conflicts are avoided. They are often used to optimize task assignments in parallel systems.
- Example: In the robotcoin problem, a conflict graph might represent scenarios where two operations cannot be performed at the same time because they access the same data in a conflicting manner.

Conflict Graph vs. Dependency Graph

Comparison in Context

- Focus:
 - The dependency graph focuses on the sequential dependencies required to compute each cell in the DP table.
 - The conflict graph focuses on which cells (tasks) cannot be updated simultaneously due to conflicts, such as resource contention.
- Types of Edges:
 - Dependency graphs have **directed edges** that show the direction of dependency.
 - Conflict graphs have **undirected edges** that indicate mutual exclusion.
- Execution Implications:
 - The dependency graph helps determine the overall **execution order** and critical path.
 - The conflict graph helps identify **parallelization constraints** by showing which operations cannot be run concurrently.

Practical application

Internal representation of compilers. ([https://ars.els-cdn.com/content/image/3-s2.0-B9780120884780000128-f12-02-9780120884780.jpg?_source:cooper torczon](https://ars.els-cdn.com/content/image/3-s2.0-B9780120884780000128-f12-02-9780120884780.jpg?_source:cooper+torczon))

Direct representation of SQL queries (https://docs.oracle.com/cd/B10500_01/server.920/a96533/scratchpad1.gif
Source: oracle)


Workflows of metagenomics analysis (<https://www.researchgate.net/profile/Ulf-Leser/publication/257799855/figure/fig5/AS:297330902355989@1447900619823/A-generic-Galaxy-workflow-for-performing-a-metagenomic-analysis-on-NGS-data-.png> source: wandelt et al. 2012) leveraging the Galaxy Framework

Project management
<https://pmatechnologies.com/wp-content/uploads/2019/09/Picture2.png> Source:
PMA Technologies

Metrics: Work

Work


- Total amount of work that is to perform on the application.
- Simply the sum of all processing times.
- Often denoted $\sum p_i$



Metrics: Work

Work

- Total amount of work that is to perform on the application.
- Simply the sum of all processing times.
- Often denoted $\sum p_i$



Here $\sum p_i = 107$


Metrics: Work

Work

- Total amount of work that is to perform on the application.
- Simply the sum of all processing times.
- Often denoted $\sum p_i$

Usage

- On m processors, the application can not be processed faster than $\frac{\sum p_i}{m}$.
- $\frac{\sum p_i}{m}$ is a **lower bound** of the **makespan**.
- $C_{max} \geq \frac{\sum p_i}{m}$




Here $\sum p_i = 107$

Metrics: Width

Width


- Maximum number of tasks that do not have direct dependencies, or transitive dependencies.
- Maximum number of independent tasks.
- Sometimes called the longest antichain.



Metrics: Width

Width

- Maximum number of tasks that do not have direct dependencies, or transitive dependencies.
- Maximum number of independent tasks.
- Sometimes called the longest antichain.



Here the width is 6.


Metrics: Width

Width

- Maximum number of tasks that do not have direct dependencies, or transitive dependencies.
- Maximum number of independent tasks.
- Sometimes called the longest antichain.

Usage

- Maximum number of useful processors.
- $\forall m > \text{Width}, S(m) = S(\text{Width})$



Here the width is 6.


Metrics: Width

Width

- Maximum number of tasks that do not have direct dependencies, or transitive dependencies.
- Maximum number of independent tasks.
- Sometimes called the longest antichain.

Usage

- Maximum number of useful processors.
- $\forall m > \text{Width}, S(m) = S(\text{Width})$



Here the width is 6.
How to find ?


Metrics: Width

Width

- Maximum number of tasks that do not have direct dependencies, or transitive dependencies.
- Maximum number of independent tasks.
- Sometimes called the longest antichain.

Usage

- Maximum number of useful processors.
- $\forall m > Width, S(m) = S(Width)$




Dilworth's algorithm

- Build a bipartite graph from dependencies
- Compute a matching
- Extract longest antichain from matching

Metrics: Critical Path

Critical Path


- Longest chain of dependency (in term of processing time)
- The length of the chain is often denoted CP , or T_{∞} .



Metrics: Critical Path

Critical Path

- Longest chain of dependency (in term of processing time)
- The length of the chain is often denoted CP , or T_{∞} .



Here $7 \rightarrow 8 \rightarrow 17 \rightarrow 16 \rightarrow 18$.
 $CP = 73$


Metrics: Critical Path

Critical Path

- Longest chain of dependency (in term of processing time)
- The length of the chain is often denoted CP , or T_{∞} .

Usage

- Whichever way the algorithm unfolds, the critical path will have to be done.
- The length of the critical path is a lower bound to the makespan
- $C_{max} \geq CP$



Here $7 \rightarrow 8 \rightarrow 17 \rightarrow 16 \rightarrow 18$.
 $CP = 73$


Metrics: Critical Path

Critical Path

- Longest chain of dependency (in term of processing time)
- The length of the chain is often denoted CP , or T_{∞} .

Usage

- Whichever way the algorithm unfolds, the critical path will have to be done.
- The length of the critical path is a lower bound to the makespan
- $C_{max} \geq CP$



Here $7 \rightarrow 8 \rightarrow 17 \rightarrow 16 \rightarrow 18$.
 $CP = 73$


Metrics: Critical Path

Critical Path

- Longest chain of dependency (in term of processing time)
- The length of the chain is often denoted CP , or T_{∞} .

Usage

- Whichever way the algorithm unfolds, the critical path will have to be done.
- The length of the critical path is a lower bound to the makespan
- $C_{max} \geq CP$



Here $7 \rightarrow 8 \rightarrow 17 \rightarrow 16 \rightarrow 18$.

$CP = 73$

How to find it?

Recursively, from the roots down

Algorithm to compute the critical path

An algorithm

Let's assume we have a Directed Acyclic Graph (DAG) representing tasks where:

- Each node v represent a task.
 - Each directed edge $u \rightarrow v$ indicates that task v depends on task u .
 - Each node v has an associated duration (processing time).
- 1 Topological Sort (DFS/Kahn's algo, checkout [here](#) for more info):
gives a linear ordering of the tasks where for every directed edge $u \rightarrow v$, task u comes before task v .
 - 2 Initialize Distances:
create a *distance* array, where $distance[v]$ stores the longest path from the start node to v .
 - 3 Calculate the distance for each node, following the order given by topological sort.
 $distance[v] = \max(distance[v], distance[u] + d(u));$
 - 4 Determine the Critical Path:
maximum value in the *distance* array.

External

Textbook:

- Chapter 2 to 5.1 of Oliver Sinnen. Task Scheduling for Parallel Systems. John Wiley & Sons, Inc. 2007. Access it through the library: <https://librarylink.uncc.edu/login?url=https://onlinelibrary.wiley.com/doi/book/10.1002/0470121173>

Cilk on graphs metrics:

- The Cilkview Scalability Analyzer, SPAA 2010. <http://web.mit.edu/willtor/www/res/cilkview-spaa-10.pdf>. a paper describing parallel application as a DAG and metrics.

Width:

- Dilworth's algorithm https://en.wikipedia.org/wiki/Dilworth%27s_theorem

Conflict graph and coloring:

- Conflict graphs: <http://math.cmu.edu/~bkell/21110-2010s/conflict-graphs.html>
- A. H Gebremedhin, F. Manne, Alex Pothén. What Color Is Your Jacobian? Graph Coloring for Computing Derivatives. Siam Review 2005.
- M. Deveci, E. Boman, K. Devine, and S. Rajamanickam. Parallel Graph Coloring for Manycore Architectures. IPDPS 2016.

Scheduling:

- A taxonomy of scheduling problems: Srishti Srivastava and Ioana Banicescu. Scheduling in Parallel and Distributed Computing Systems. Chapter 11 of Prasad, Gupta, Rosenberg, Sussman, and Weems. Topics in Parallel and Distributed Computing: Enhancing the Undergraduate Curriculum: Performance, Concurrency, and Programming on Modern Platforms, Springer International Publishing, 2018. https://grid.cs.gsu.edu/~tcpp/curriculum/?q=system/files/Ch11_4.pdf
- Scheduling is NP-Hard: M. Garey and D. Johnson. Computers and Intractability: A Guide to the Theory of NP-Completeness. Freeman. 1979.
- LS for independent tasks: R. Graham. Bounds for certain multiprocessing anomalies. Bell System Technical Journal. 1966
- LPT and LS with precedence: R. Graham. Bounds on Multiprocessing Timing Anomalies. SIAM Journal on Applied Mathematics. 1969.
- Chapter 1 and 7 of. H. Casanova, A. Legrand, Y. Robert. Parallel Algorithms, CRC Press. 2008